

Physics-Aware Multi-Objective Scheduling for an Agile SAR Satellite: Balancing Profit, NESZ, and Geometric Resolution

LIU Shuhao

College of Intelligent Systems Science and Engineering, Harbin Engineering University, 145 Nantong Street, Nangang District, Harbin, 150001, China

Abstract

In agile SAR scheduling, observation timing determines imaging quality through radar-specific geometric and radiometric effects. We characterize an empirical saturation regime for multi-objective SAR quality optimization in single-pass, single-satellite Stripmap operation by integrating Noise-Equivalent Sigma Zero (NESZ, radiometric) and geometric resolution as decomposed objectives within an NSGA-III framework. Across densities 20–500, saturation is scale-dependent: at $N = 20$, MOEA-2 extracts a +4.5% gain in geometric quality f_2 at a 13% coverage cost, and adding the NESZ objective f_3 yields a further +58% radiometric gain over the greedy baseline (G-BL) at 2% coverage cost (and $\sim 6\%$ geometric-quality cost) for MOEA-3; at operational densities ($N \geq 300$), multi-objective evolutionary algorithm (MOEA) formulations seeded by G-BL approach the coverage anchor and the marginal quality gain falls below 2%. This convergence is conditional on G-BL hot-start initialization, with random-init f_1^* deficits of 0.24–0.62 peaking at intermediate densities; ablation shows removing the physical objectives barely changes dense-regime outcomes. The third objective is therefore density-dependent rather than redundant: MOEA-3 differs markedly from MOEA-2 at sparse density (+84% f_3) and converges at operational densities, reflecting lower-squint geometric coupling rather than a mathematical identity. Rather than identifying a superior solver, our results establish an empirical saturation boundary for physics-aware multi-objective scheduling.

Email address: liushuhao@hrbeu.edu.cn (LIU Shuhao)

Keywords: agile SAR, satellite scheduling, multi-objective optimization, NSGA-III, noise-equivalent sigma zero (NESZ), geometric resolution

1. Introduction

Agile Earth-observation satellites equipped with Synthetic Aperture Radar (SAR) are increasingly central to time-critical monitoring missions, from disaster response to maritime surveillance. Their agility (the ability to slew the antenna beam away from the nadir track) enables a single satellite to image far more targets per orbit than a fixed-pointing platform (Curlander and McDonough, 1991). This flexibility, however, transforms scheduling from a simple feasibility check into a combinatorial optimization problem: for N targets each with a finite visibility window, the scheduler must select a subset and assign observation times that satisfy attitude transition constraints while maximizing some measure of mission value. In conventional satellite scheduling, observation quality is commonly represented as a fixed attribute of each pre-computed opportunity rather than a quantity that varies with the selected acquisition time. The scheduling problem therefore reduces to maximizing the number or priority-weighted sum of observed targets subject to temporal feasibility. This assumption underpins most existing satellite scheduling work, from mixed-integer programming formulations to deep reinforcement learning approaches (Chen et al., 2025).

SAR imaging inherently departs from this assumption because image quality depends on the radar geometry at the moment of acquisition: the ground-range resolution degrades with incidence angle ($\rho_{\text{rg}} \propto 1/\sin \theta$), the azimuth resolution broadens with squint angle ($\rho_{\text{az}} \propto 1/\cos \psi_{\text{sq}}$), and the radiometric sensitivity (Noise-Equivalent Sigma Zero, NESZ) deteriorates with slant range as R^3 (Cumming and Wong, 2005). When the scheduler can choose *when* within a visibility window to observe each target (the defining feature of agile satellites), it simultaneously chooses the imaging quality at which each target is acquired. Despite this tight coupling between scheduling and imaging physics, with few exceptions (Wei et al., 2025; Kramer and Miller, 2025), existing agile SAR scheduling methods either reduce SAR to a coverage-only objective or apply post hoc quality checks on schedules produced by coverage-maximizing solvers. Few approaches treat physical quality dimensions as primary optimization objectives in the SAR domain.

The problem is NP-hard by reduction from the Orienteering Problem with Time-Dependent Profits (Peng et al., 2019), precluding polynomial-

time exact algorithms for all but the smallest instances. The SAR literature has well-established models for NESZ and geometric resolution as functions of observation geometry (Curlander and McDonough, 1991; Cumming and Wong, 2005), and the multi-objective optimization literature provides mature frameworks such as the Non-dominated Sorting Genetic Algorithm III (NSGA-III) (Deb and Jain, 2014) for handling three or more objectives simultaneously. Yet direct integration of these three elements (SAR physical modeling, multi-objective evolutionary algorithms (MOEAs), and satellite scheduling) remains limited. We therefore formulate SAR scheduling with separate, physics-derived geometric and radiometric quality objectives within a MOEA framework.

We formulate a three-objective agile SAR scheduling problem in which the scheduler optimizes (1) total coverage profit, (2) mean geometric image quality (ground-range and azimuth resolution), and (3) mean NESZ radiometric quality simultaneously. We compare five solvers spanning the methodological spectrum: two greedy heuristics (one profit-only baseline and one squint-minimized variant), a single-objective genetic algorithm with greedy hot-start, and two NSGA-III variants with two and three objectives. The solvers are evaluated on 200 systematically varied scenarios across four density classes ($N = 20$ to $N = 500$). Our results show a scale-dependent saturation, not a uniform solver advantage. In sparse scenarios, multi-objective optimization can extract substantial radiometric gains at little coverage cost (MOEA-3 raises f_3 by +58% relative to G-BL at only 2% coverage cost) alongside modest geometric gains (MOEA-2, +4.5% in f_2 at 13% coverage cost). In dense scenarios, the multi-objective advantage diminishes to negligible levels: greedy methods attain the tested family’s coverage ceiling, and the squint-minimized heuristic offers competitive radiometric quality. Section 7.1 provides a detailed synthesis of these findings. The specific contributions of the paper are:

1. The discovery and analytical characterization of an empirical saturation regime for multi-objective SAR quality optimization in the single-pass, single-satellite Stripmap setting: the regime arises from solver convergence to lower-squint geometries that attenuate the envelope-level conflict between geometric and radiometric quality while coverage pressure removes the trading freedom between them (§ 6.4), not from a physical identity between objectives, and renders the three-objective formulation redundant at operational densities. A controlled ablation study (four objective-function variants, 200 scenarios per variant) con-

firmly that solver outcomes become invariant to the choice of physical objectives at high target densities, and initialization controls (§ 6.6) confirm that coverage convergence is conditional on greedy-baseline (G-BL) hot-start initialization;

2. The decomposition of SAR image quality into closed-form NESZ ($\propto R^3$) and geometric quality index ($\sin \theta \cos \psi_{sq}$) objectives within an NSGA-III framework, distinct from the scalar visibility composite of Kramer and Miller (2025) and the generic image-quality composite of Wei et al. (2025), using mean per-task quality metrics that decouple quality from the number of scheduled tasks. This formulation enables the systematic characterization of quality–coverage trade-offs across five solvers spanning greedy heuristics, single-objective GA, and NSGA-III variants on four scale classes (200 scenarios).

2. Related Work

We examine three lines of research that our study connects: satellite scheduling (2.1), SAR physical quality modeling (2.2), and multi-objective evolutionary optimization (2.3).

2.1. Satellite Scheduling

Satellite scheduling has been studied extensively in the Earth-observation domain. The survey by Ferrari et al. (2025) catalogs over 200 publications spanning exact methods, heuristics, metaheuristics, and learning-based approaches. Exact formulations based on mixed-integer linear programming (MILP) achieve optimality on small instances but fail to scale to operational problem sizes, motivating a broad family of heuristic and metaheuristic alternatives including genetic algorithms, simulated annealing, and tabu search (Wang et al., 2019; Chen et al., 2019). The earliest formal treatments of agile satellite scheduling by Lemaître et al. (2002) and Bianchessi et al. (2007) established the foundational problem structure of time-dependent transition times between observations. Recent work further generalizes these methods to agile satellites, where the freedom to choose observation start times within each visibility window is exploited to pack more observations into a single orbit (Peng et al., 2019; Chen et al., 2024). A parallel line of research has pursued learning-based and hybrid approaches for satellite scheduling. Liu et al. (2024) and Zhou et al. (2025) apply deep reinforcement learning to agile optical satellite scheduling, while Jacquet et al. (2024)

combine graph neural networks with Monte Carlo tree search for the same domain. Wang et al. (2026) propose a Lagrangian relaxation-based heuristic for multi-satellite mission planning. Caleiras et al. (2026) formulate super-agile satellite scheduling as a constraint programming problem, capturing the time-dependent transition constraints characteristic of super-agile platforms. Herrmann and Schaub (2023) apply RL to on-board AEOS scheduling, maximizing target collection reward without imaging quality as an objective. Wu et al. (2026) address high-target-density AEOS scheduling with an improved whale optimization algorithm, reporting convergence challenges that echo the density-driven constraint tightening we analyze in § 6.3. Despite their methodological diversity, most prior works treat quality as a fixed property determined at observation-window generation. In contrast, SAR imaging quality varies continuously with acquisition geometry and is directly optimizable within the scheduling window. Chang et al. (2022) propose early MOEAs for SAR satellite scheduling that combine an adaptive large neighborhood search with NSGA-II and NSGA-III selection; their three-objective model (ground-target loss rate, invalid observations, and sensor-activation count) targets operational metrics — the closest prior work to ours — but it omits imaging quality dimensions. Wei et al. (2025) introduce a neural-policy-guided MOEA that incorporates a composite image quality score as an optimization objective. Kramer and Miller (2025) schedule SAR constellation observations with a hybrid metaheuristic using a weighted visibility-performance composite, embedding a SAR-quality proxy in the scheduler, and Yao et al. (2024) develop a bilevel evolutionary algorithm for multi-satellite scheduling that separates task allocation from detailed observation planning. Chang and Zhou (2025) formulate agile Earth-observation scheduling with a bi-objective image-quality-loss/energy model in which quality enters as an attitude-angle proxy attached to each observation opportunity, and Wu et al. (2019) apply NSGA-III to satellite-formation imaging scheduling with coverage- and timeliness-type objectives; both are representative of Chinese-language and Chinese-affiliated scheduling research in which imaging quality is a scalar attribute or coarse proxy rather than a decomposed physical objective. The common limitation of this prior work is that SAR physical quality metrics—noise-equivalent sigma zero (NESZ), azimuth resolution, ground-range resolution—are either absent or approximated by coarse proxies (e.g., penalizing extreme off-nadir angles). To our knowledge, no existing SAR scheduling method treats continuous, physics-derived quality functions as decomposed optimization objectives in their own right: the closest ap-

proaches embed aggregated quality proxies rather than separate geometric and radiometric objectives.

2.2. SAR Physical Quality Modeling

The SAR remote sensing literature provides well-established analytical models linking observation geometry to image quality. Curlander and McDonough (1991), Cumming and Wong (2005), and Moreira et al. (2026) give detailed treatments of SAR physics, including the two quality dimensions relevant to scheduling: geometric resolution and radiometric sensitivity. Geometric resolution decomposes into ground-range resolution $\rho_{\text{rg}} = c/(2B \sin \theta)$, where c is the speed of light, B the radar bandwidth, and θ the incidence angle, and azimuth resolution $\rho_{\text{az}} = D/(2 \cos \psi_{\text{sq}})$, where D is the antenna length and ψ_{sq} the squint angle. In operational SAR, amplitude weighting widens ρ_{az} by 10–20% (Curlander and McDonough, 1991), but since this is platform-constant, the unweighted form preserves the relative ordering for optimization. The two resolution components respond oppositely to departure from the zero-Doppler geometry ($\psi_{\text{sq}} = 0$, θ minimal): ρ_{az} worsens as $|\psi_{\text{sq}}|$ grows, whereas ρ_{rg} is coarsest at the minimal incidence of closest approach and improves as θ rises away from it—the origin of the θ -axis trade-off formalized in § 3.2. Radiometric sensitivity, measured by the Noise-Equivalent Sigma Zero (NESZ), deteriorates with the cube of slant range: $\text{NESZ} \propto R^3$ for Stripmap SAR after azimuth compression gain (Curlander and McDonough, 1991; Moreira et al., 2026). We adopt a 3D geometric approximation that decomposes the slant range as $R \approx R_{\text{elev}}/\cos \psi_{\text{sq}}$ (where R_{elev} is the elevation-plane range at zero squint), yielding $\text{NESZ} \propto 1/(\cos^3 \theta \cdot \cos^3 \psi_{\text{sq}})$. This plane-Earth range model omits the Earth’s curvature and the antenna elevation-pattern roll-off; at the 47° edge of the tested incidence range it overestimates slant range by roughly 5% and R^3 by roughly 17%, and it suppresses the $\sin \theta$ factor that a ground-resolution reference area adds to standard NESZ expressions. It is therefore used here as a relative ordering index across observation geometries, not as an absolute radiometric prediction – the comparisons in §5.3 and below depend only on the ordering it induces. The $\cos^3 \psi_{\text{sq}}$ factor enters through the slant-range decomposition. In the schedules produced by the solvers below, the per-task squint angle is moderate rather than near-zero (the mean $|\psi_{\text{sq}}|$ across the four density classes is $16\text{--}24^\circ$, with a 90th percentile of $32\text{--}39^\circ$), so $\cos^3 \psi_{\text{sq}}$ contributes on the order of 10–25% of the f_3 range rather than the few percent it would at near-zero squint. The proportionality $\text{NESZ} \propto R^3$ captures

the dominant geometric dependence; platform-constant terms (noise figure, transmit power, processing gain) do not affect the relative ordering of observation geometries. This cubic dependence means even modest deviations from the optimal viewing geometry make fine observation timing critical for radiometric performance. These quality models are routinely used in SAR system design but their use in scheduling remains limited to post hoc validation of fixed observation plans (Kramer and Miller, 2024) or binary feasibility filtering at the observation-window level (Li et al., 2026). The scheduling literature and the SAR quality modeling literature have developed in parallel: schedulers treat quality as a constraint, while SAR engineers compute quality after the schedule is fixed. Our study connects these two domains by embedding continuous quality functions directly into the optimization objective space.

2.3. Multi-Objective Evolutionary Optimization

Multi-objective evolutionary algorithms (MOEAs) are a standard toolkit for problems with competing objectives. The NSGA-III algorithm (Deb and Jain, 2014) extends the NSGA-II framework to many-objective problems through reference-point-based diversity preservation. The hypervolume (HV) indicator is widely used for comparing Pareto front approximations, but its sensitivity to frontier cardinality creates interpretational challenges when mixing single-objective and multi-objective solvers, a methodological issue we address in § 6.4. To our knowledge, MOEAs have not been applied to SAR scheduling with physical quality objectives; the closest work uses single-objective optimization with post hoc quality evaluation. Our NSGA-III formulation makes these trade-offs explicit by placing coverage profit, geometric quality, and radiometric quality in a unified multi-objective framework.

3. Problem Formulation

3.1. System Model

We consider a single agile SAR satellite in a circular orbit at altitude H , operating in Stripmap mode, which provides the widest continuous swath with uniform quality and analytically tractable NESZ and resolution expressions (Curlander and McDonough, 1991). A problem instance $\mathcal{I} = (\mathcal{T}, \mathcal{P}, \mathcal{O})$ consists of a set of N ground targets $\mathcal{T}$, a satellite platform $\mathcal{P}$, and a SAR imaging model $\mathcal{O}$.

Task Set. Each task $i \in \{1, \dots, N\}$ is specified by its geographic coordinates $\mathbf{r}_i = (\phi_i, \lambda_i)$, a priority weight $p_i \in \mathbb{R}^+$, a minimum required ground resolution ρ_i^{req} , a visibility window $\mathcal{W}_i = [a_i, b_i]$ during which the satellite can observe the target, and a fixed observation duration d_i .

Satellite Platform. The satellite is characterized by its maximum slew rate ω_{max} (a conservative single-axis bound on three-axis attitude maneuvers), settling time τ_{settle} after each maneuver, per-orbit energy budget E_{max} and memory budget M_{max} , and per-observation energy e_i and memory m_i consumption.

SAR Observables. Given an observation time t_i , the satellite-to-target line-of-sight (LOS) vector $\mathbf{l}(t_i)$ in the Local-Vertical Local-Horizontal (LVLH) frame is fully determined by the satellite’s orbital state. From $\mathbf{l}(t_i)$ we derive four geometric quantities deterministically: slant range $R = \|\mathbf{l}\|$, off-nadir angle $\xi = \arctan 2(\sqrt{l_x^2 + l_y^2}, l_z)$, incidence angle θ (via the Earth-curvature-corrected mapping $\theta = \arcsin(\frac{R_e + H}{R_e} \sin \xi)$ where R_e is Earth’s mean radius), and squint angle $\psi_{\text{sq}} = \arcsin(|l_x|/R)$ (LVLH x -axis along-track, z -axis nadir-pointing), defined as the angular deviation of the LOS vector from the zero-Doppler plane (Curlander and McDonough, 1991). A critical property is that *all* quality-relevant geometry is a deterministic function of the single decision t_i .

Decision Variables. A schedule consists of a task selection vector $\mathbf{x} \in \{0, 1\}^N$, an ordered sequence $S = (s_1, \dots, s_m)$ of the m selected tasks, and a vector of observation start times $\mathbf{t} = (t_{s_1}, \dots, t_{s_m})$. All geometric parameters ($\theta_i, \psi_{\text{sq},i}, R_i$) are *derived* from t_i rather than independent decision variables.

3.2. Objective Functions

We formulate a three-objective optimization problem that captures the tension between scheduling throughput and SAR-specific imaging quality. All quality objectives use the *mean* over scheduled tasks, making them independent of the number of scheduled tasks $m = |\mathcal{S}|$ and thereby producing a genuine “quantity versus quality” Pareto trade-off.

Objective 1: Coverage Profit. The normalized coverage profit (Eq. 1) is defined relative to the greedy baseline

$$f_1^* = \frac{1}{f_1^{\text{G-BL}}} \sum_{i \in \mathcal{S}} p_i \quad (1)$$

where $f_1^{\text{G-BL}}$ is the profit achieved by the profit-only greedy scheduler (G-BL, § 4.2) on the same scenario. Normalization by G-BL rather than GA-P-BL is conservative for the MOEA coverage deficit: GA-P-BL output is identical to G-BL ($f_1^* = 1.00$ at every scale, Table 3), so G-BL-normalized f_1^* values equal the coverage comparison to the single-objective GA.

Objective 2: Geometric Image Quality.. The ground pixel area of a SAR image is the product of ground-range resolution $\rho_{\text{rg}} = c/(2B \sin \theta)$ and azimuth resolution $\rho_{\text{az}} = D/(2 \cos \psi_{\text{sq}})$. We define a dimensionless geometric quality index:

$$f_2 = \frac{1}{|\mathcal{S}|} \sum_{i \in \mathcal{S}} \sin \theta_i \cdot \cos \psi_{\text{sq},i} \quad (2)$$

Higher f_2 indicates smaller ground pixel area (better geometric resolution). The quality indices are unweighted means over the scheduled set, orthogonal to the priority weights used in f_1 : they reflect the average imaging geometry of the selected tasks rather than biasing quality toward high-priority targets.

Objective 3: NESZ Radiometric Quality.. Noise-Equivalent Sigma Zero (NESZ) measures the minimum detectable backscatter. For Stripmap SAR, NESZ degrades with the cube of slant range, $\text{NESZ} \propto R^3 \propto 1/(\cos^3 \theta \cdot \cos^3 \psi_{\text{sq}})$, after accounting for azimuth compression gain (Curlander and McDonough, 1991). We define the radiometric quality index as:

$$f_3 = \frac{1}{|\mathcal{S}|} \sum_{i \in \mathcal{S}} \cos^3 \theta_i \cdot \cos^3 \psi_{\text{sq},i} \quad (3)$$

Higher f_3 indicates lower NESZ (better radiometric sensitivity).

Physical Trade-Off.. The two quality objectives partition SAR imaging performance by physical mechanism: f_2 captures spatial resolution (geometric), f_3 captures signal-to-noise ratio (radiometric). Their optima are not coincident; f_3 peaks at the zero-Doppler point (θ minimal, $\psi_{\text{sq}} = 0$), whereas f_2 benefits from moderately higher θ because $\sin \theta$ grows with incidence angle. This creates a genuine trade-off: pursuing the cleanest NESZ (zero-Doppler observation) forces acceptance of stretched ground-range pixels; pursuing optimal ground pixel area requires deviating from the zero-Doppler point and accepting elevated radiometric noise. The squint angle ψ_{sq} appears in both f_2 and f_3 (Eqs. 2–3): reducing ψ_{sq} simultaneously improves both resolution components through the $\cos \psi_{\text{sq}}$ factor. In practice, the zero-squint

point coincides with the time of closest approach, which minimizes the slant range and, for a spherical Earth, approximately minimizes the incidence angle θ and thus $\sin \theta$; reducing ψ_{sq} therefore lowers f_2 through a geometric (not algebraic) coupling between the two angles, creating a two-axis $[\theta, \psi_{\text{sq}}]$ trade-off structure that limits multi-objective gain once solvers converge to lower-squint regimes (Eq. 6 in § 6.4 and its axis decomposition for the separation into definitional and scheduler-induced components).

Because θ and ψ_{sq} are derived from the same LOS vector (§ 3.1), the per-task f_2 and f_3 contributions share geometric structure; its quantitative characterization is presented in § 6.4.

3.3. Constraints

Observation-window membership is treated as the domain of each start-time variable rather than as a separate numbered constraint:

$$t_i \in [a_i, b_i - d_i] \quad \forall i \in \mathcal{S}. \quad (4)$$

The remaining four constraints couple task selection, acquisition geometry, temporal feasibility, and resource use. The incidence-angle and squint-angle constraints (C1) are enforced during visibility-window generation; C2–C4 are checked inline by every solver.

C1: Incidence and Squint Angle Constraints.. Each selected acquisition must satisfy the platform incidence-angle range, $\theta^{\min} \leq |\theta_i| \leq \theta^{\max}$, and the squint-angle limit $|\psi_{\text{sq},i}| \leq \psi_{\text{sq}}^{\max}$, where $\psi_{\text{sq}}^{\max}$ is the satellite’s maximum achievable squint (assumed $\leq 45^\circ$; real agile SAR platforms typically slew only a few degrees, so this is a conservative envelope for the tested agility). Both are observation-geometry limits applied before scheduling.

C2: Attitude Maneuver and Non-Overlap.. Between consecutive observations, the satellite must rotate through the angular separation of their LOS vectors. Let $\mathbf{l}_k$ and $\mathbf{l}_{k+1}$ be the LVLH-frame LOS vectors at t_{s_k} and $t_{s_{k+1}}$. The angular gap $\Delta\eta_k = \arccos(\mathbf{l}_k \cdot \mathbf{l}_{k+1} / (|\mathbf{l}_k||\mathbf{l}_{k+1}|))$ must be traversed within the available inter-observation time:

$$t_{s_{k+1}} \geq t_{s_k} + d_{s_k} + \frac{\Delta\eta_k}{\omega_{\max}} + \tau_{\text{settle}} \quad \forall k \in \{1, \dots, m-1\}. \quad (5)$$

This inequality also enforces non-overlapping task-execution intervals. It is the core coupling in our formulation: changing t_{s_k} shifts $\mathbf{l}_k$, affecting both imaging quality (via θ_k and $\psi_{\text{sq},k}$) and transition feasibility (via $\Delta\eta_k$).

The tightness of C2 scales with task density: as N grows, the fraction of task pairs whose required angular transition $\Delta\eta/\omega_{\max}$ exceeds the available inter-observation time increases, tightening the constraint and making C2 the dominant active constraint in dense regimes.

C3–C4: Energy and Memory Budgets. $\sum_{i \in \mathcal{S}} e_i \leq E_{\max}$ and $\sum_{i \in \mathcal{S}} m_i \leq M_{\max}$, where $E_{\max} = 10^7$ and $M_{\max} = 10^{11}$ (arbitrary units) are set sufficiently high to be non-binding under all tested scenarios (typical per-task $e_i \sim 10^2\text{--}10^3$, $m_i \sim 10^6\text{--}10^7$, far below the caps); they are retained for completeness.

4. Methodology

4.1. Solver Overview

Table 1 summarizes the five solvers compared in this study. We describe the solvers in order of increasing sophistication: from greedy heuristics (G-BL, G-SM) through single-objective metaheuristic (GA-P-BL), to multi-objective evolutionary algorithms (MOEA-2, MOEA-3).

Table 1: Solvers: greedy baselines (G-BL, G-SM), single-objective GA with hot-start (GA-P-BL), and NSGA-III variants with two and three objectives (MOEA-2, MOEA-3)

Label	Type	Objectives	Hot-Start
G-BL	Greedy heuristic	$\max f_1$ (profit-only)	n/a
G-SM	Greedy heuristic	$\max f_1 + \text{squint minimization}$	n/a
GA-P-BL	Single-objective GA	$\max f_1$ (G-BL seeded)	std = 0.5
MOEA-2	NSGA-III (2-objective)	$f_1^* + f_2$ (geometric quality)	std = 0.5
MOEA-3	NSGA-III (3-objective)	$f_1^* + f_2 + f_3$ (geo.+NESZ)	std = 0.5

Single-objective solvers (G-BL, G-SM, GA-P-BL) optimize only f_1 ; their f_2 and f_3 values are computed post hoc on the final schedule for comparison with the multi-objective solvers. The five solvers differ primarily in objective-function configuration rather than algorithmic paradigm. This design is intentional: it isolates the effect of adding physical quality objectives, which is the paper’s research question. A comparison across algorithmic families (e.g., NSGA-III vs. MOEA/D (decomposition-based) vs. SMS-EMOA (hypervolume-based)) is left to future work (§ 7.3).

4.2. G-BL: Greedy Baseline

G-BL represents the conventional scheduling approach: it makes no use of SAR imaging physics when ordering or selecting tasks. Tasks are sorted by earliest start time t_{start} and greedily inserted into the schedule at their earliest feasible time. This earliest-time placement fixes each task’s observation geometry implicitly through the window’s incidence–squint profile rather than through any quality criterion. If the attitude transition constraint (C2) is violated for a candidate task, its start time is delayed within the window until the constraint is satisfied; if no feasible delay exists, the task is dropped. Since G-BL selects observation times solely for feasibility, its f_2 and f_3 values are incidental by-products of the window geometry. We compute these post hoc on the final schedule to serve as the baseline against which physics-aware solvers are compared.

4.3. G-SM: Greedy Squint-Minimized

G-SM shares G-BL’s greedy structure (sort by earliest start time) but modifies the observation time selection: within each visibility window, a task is placed at the discretized time step that minimizes squint angle $\psi_{\text{sq}} = \arcsin(|l_x|/R)$. G-SM minimizes squint angle as a computationally cheap proxy for maximizing f_3 (NESZ radiometric quality); the correlation between squint minimization and f_3 improvement is confirmed by the large effect size (Cliff’s $\delta = +1.0$, § 6.2). This strategy serves as a physics-aware greedy baseline: it answers the question “what quality improvement can a simple squint-minimization heuristic achieve, without any multi-objective search?” The squint-minimizing time selection incurs a profit penalty because the optimal-squint time may lie later in the window, reducing the remaining time for subsequent tasks (constraint C2). As with G-BL, f_2 and f_3 are post hoc outputs.

4.4. GA-P-BL: Single-Objective GA with Hot-Start

GA-P-BL tests whether a genetic algorithm optimizing only f_1 can outperform greedy scheduling by exploiting temporal flexibility. Each chromosome encodes a vector of normalized start times $\boldsymbol{\tau} \in [0, 1]^N$, where task i ’s actual start time is $t_i = t_i^{\text{earliest}} + \tau_i \cdot (t_i^{\text{latest}} - d_i - t_i^{\text{earliest}})$, alongside a binary selection vector $\mathbf{x} \in \{0, 1\}^N$; both are decision variables (a $2N$ -dimensional chromosome). The population of 100 individuals is initialized from G-BL’s solution with Gaussian noise (std = 0.5) added to each τ_i . The GA runs for 200 generations with simulated binary crossover ($p_c = 0.9$, $\eta_c = 20$) and

polynomial mutation ($p_m = 0.1$, $\eta_m = 20$). f_2 and f_3 are computed post hoc on the best-found schedule. After optimization, a never-worse-than-baseline guard compares the best GA schedule with G-BL: if the GA’s f_1 falls short, the greedy schedule is returned instead, guaranteeing $f_1^{\text{GA-P-BL}} \geq f_1^{\text{G-BL}}$ by construction. In practice this guard makes GA-P-BL output numerically identical to G-BL on every scenario (§6.1); GA-P-BL therefore acts as a control that confirms the greedy profit baseline rather than as an independent solver.

4.5. MOEA-2 and MOEA-3: Multi-Objective Evolutionary Algorithms

MOEA-2 and MOEA-3 extend the GA-P-BL chromosome encoding to multi-objective optimization via the NSGA-III framework (Deb and Jain, 2014). Both use a population of 100 individuals over 200 generations with the same genetic operators as GA-P-BL. Reference points are generated via the Das–Dennis systematic approach, producing $H = \binom{n_{\text{obj}}+p-1}{p}$ structured reference directions, where n_{obj} is the number of objectives and $p = 12$ the divisions per objective dimension ($H = 13$ for MOEA-2, $H = 91$ for MOEA-3). The entire population is hot-started from the G-BL solution perturbed by Gaussian noise ($\text{std} = 0.5$), and an at-least-one-task constraint prevents empty solutions from dominating the Pareto front. As every solver is hot-started from the non-empty G-BL solution (§ 4.2), this constraint is inactive in the reported experiments. MOEA-2 optimizes two objectives: f_1^* (normalized profit) and f_2 (mean geometric quality, Eq. 2). MOEA-3 adds f_3 (mean NESZ radiometric quality, Eq. 3) as a third objective. The comparison between MOEA-2 and MOEA-3 directly tests whether adding the NESZ objective (f_3) produces Pareto-front differentiation beyond what f_2 alone captures, or whether the θ/ψ_{sq} geometric coupling renders the third objective redundant in the single-pass, single-satellite setting.

5. Experimental Setup

5.1. Scenario Design

Table 2 defines the four scenario classes.

Table 2: Scenario classes: four density levels (S1–S4) with $N = 20, 100, 300, 500$ targets, 50 scenarios per class

Class	Targets	Look Direction	Description
S1	20	Both	Small, low-density
S2	100	Both	Medium, moderate density
S3	300	Both	Medium-large, transition region
S4	500	Both	Large, high-density

All scenarios share a common satellite platform abstracted as a generic C-band Stripmap SAR with bilateral-looking capability. The platform follows a circular sun-synchronous orbit at 693 km altitude (inclination $\approx 97.8^\circ$, LTAN 06:00, dawn–dusk) and carries a C-band Stripmap instrument with antenna length $D = 12.3$ m and bandwidth $B = 56.5$ MHz. It accesses both left- and right-side incidence angles in the range $[18^\circ, 47^\circ]$ (signed convention: $\theta \in [-47^\circ, -18^\circ] \cup [18^\circ, 47^\circ]$). This generic bilateral configuration is selected to exercise the full incidence–squint trade-off space.

Targets are distributed along the satellite ground track within the accessible swath (cross-track half-width $\pm 5.5^\circ$ from nadir, roughly ± 608 km). Within each class, the five scenario variants (labeled A–E in the scenario files) span the target distribution (uniform, clustered, and mixed), and three variants additionally modulate the geometry rather than only the distribution: S1-D and S3-D use a single orbital revolution (yielding fewer and narrower windows), S2-E places targets at high latitudes (± 40 – 70°), and S4-D narrows the incidence ceiling from 47° to 25° . The 50 scenarios per class consist of these 5 variants $\times$ 10 random seeds each; the density trend reported in §6.3 is robust to these within-class modulations (e.g., the sparse-regime coverage deficit appears across all S1 variants, and the densest S4 class remains at or near full coverage even for the variants with fewer or tighter windows). The four density classes (S1–S4) span sparse reconnaissance ($N = 20$) to dense monitoring ($N = 500$). Other intermediate densities ($N = 50, 400$) are omitted to limit computational cost. To directly constrain the transition midpoint, two additional intermediate classes—S7 ($N = 150$) and S8 ($N = 200$), 50 scenarios each generated under the same generic C-band bilateral configuration—were run for MOEA-2 and the two greedy baselines (G-BL, G-SM) and are used solely in the scale-sensitivity analysis in § 6.3; they do not enter the five-solver pooled comparisons.

Each task is assigned a priority weight p_i drawn uniformly from the integer set $\{1, 2, \dots, 10\}$, and a fixed observation duration $d_i = 30$ s. Visibility windows are computed via orbit propagation over two orbital periods (~ 200 min), yielding 1–2 windows per target on average, each spanning several minutes. Windows are filtered by incidence-angle and minimum-elevation (5° above the local horizon) constraints. The resulting windows are approximately evenly split between left- and right-looking opportunities ($\sim 50\%$ each).

5.2. Solver Configuration

All five solvers run on a single CPU core to ensure a uniform computational environment. We do not report wall-clock runtimes: the greedy heuristics complete in milliseconds while the evolutionary solvers take minutes on the densest scenarios, and the study’s claims concern solution quality rather than computational efficiency. G-BL and G-SM are deterministic greedy heuristics with no tunable hyperparameters. GA-P-BL and the MOEA variants use the population size, generation budget, and genetic operators described in §4.4 and §4.5, with the entire MOEA population hot-started from the G-BL solution.

5.3. Evaluation Metrics

We use hypervolume (HV) as the primary quality indicator for multi-objective comparison. Statistical significance is assessed via the Friedman test for overall differences across 5 solvers $\times$ 200 scenarios, followed by pairwise Wilcoxon signed-rank comparisons with Bonferroni correction ($\alpha = 0.05/10 = 0.005$). Per-objective and ablation comparisons are reported with raw p -values alongside the Bonferroni threshold. Effect sizes are measured by Cliff’s δ , classified following Romano, Kromrey, Coraggio, and Skowronek (2006) ($|\delta| < 0.147$ negligible, < 0.33 small, < 0.474 medium, ≥ 0.474 large). Hypervolume is computed with reference point $(0, 0, 0)$ in the min–max-normalized objective space, where the min/max bounds are pooled across all solvers and scenarios (a single global normalization basis). This shared normalization framework makes cross-solver HV values directly comparable as relative rankings; the caveat documented in §6.4 is that HV also reflects the cardinality of each solver’s reported non-dominated front.

5.4. Experiment Design Overview

This study comprises three classes of experiments. The first is the main experiment, running all five solvers (G-BL, G-SM, GA-P-BL, MOEA-2, MOEA-3) on 50 scenarios per density class (S1–S4), totaling 200 scenarios. Results from the main experiment are presented in §6.1–§6.4 and §6.2. The second is the ablation experiment, which successively removes the squint component (variant B), the incidence-angle component (variant C), and all physical objectives (variant D) from the MOEA-3 framework, running each variant on the same 200 scenarios (§6.5). The third comprises robustness control experiments—hot-start dependence, perturbation-sensitivity sweep, random-search calibration, and search-budget control—which verify that the main and ablation findings are not artifacts of initialization or parameter choices (§6.6).

6. Results and Analysis

6.1. Overall Solver Performance

Table 3 reports all five solvers across all four scenario classes. The full-scale view shows why the value of multi-objective optimization cannot be judged from S4 alone. At the sparsest density S1, MOEA-2 improves f_2 by +4.5% relative to G-BL while accepting a 13% coverage reduction ($f_1^* = 0.87$), and MOEA-3 raises f_3 by +58% at only 2% coverage cost ($f_1^* = 0.98$). At S2 the effects contract, and at S3–S4 MOEA coverage and quality converge to the G-BL reference. Across every scale, GA-P-BL output is numerically identical to G-BL, because its never-worse-than-baseline guard (Sec. 4.4) replaces the GA schedule whenever it fails to beat the greedy anchor on the single profit objective. GA-P-BL therefore serves as a control that confirms the greedy coverage baseline rather than an independent solver. G-SM occupies a different operating regime, exchanging coverage and f_2 for substantially higher f_3 .

Table 3: Solver performance across all scenario classes (mean $\pm$ standard deviation over 50 scenarios per class). Higher values are preferred for f_1^* , f_2 , and f_3 .

Class	Solver	f_1^*	f_2	f_3
S1 ($N = 20$)	G-BL	1.00 ± 0.00	0.574 ± 0.026	0.271 ± 0.057
	G-SM	0.51 ± 0.16	0.475 ± 0.057	0.644 ± 0.058
	GA-P-BL	1.00 ± 0.00	0.574 ± 0.026	0.271 ± 0.057
	MOEA-2	0.87 ± 0.08	0.600 ± 0.021	0.233 ± 0.027
	MOEA-3	0.98 ± 0.03	0.540 ± 0.021	0.428 ± 0.064
S2 ($N = 100$)	G-BL	1.00 ± 0.00	0.549 ± 0.033	0.366 ± 0.084
	G-SM	0.35 ± 0.09	0.460 ± 0.054	0.649 ± 0.052
	GA-P-BL	1.00 ± 0.00	0.549 ± 0.033	0.366 ± 0.084
	MOEA-2	0.98 ± 0.03	0.556 ± 0.031	0.347 ± 0.077
	MOEA-3	1.00 ± 0.00	0.548 ± 0.031	0.375 ± 0.075
S3 ($N = 300$)	G-BL	1.00 ± 0.00	0.527 ± 0.020	0.452 ± 0.064
	G-SM	0.29 ± 0.04	0.437 ± 0.036	0.664 ± 0.035
	GA-P-BL	1.00 ± 0.00	0.527 ± 0.020	0.452 ± 0.064
	MOEA-2	1.00 ± 0.00	0.530 ± 0.019	0.444 ± 0.062
	MOEA-3	1.00 ± 0.00	0.528 ± 0.020	0.453 ± 0.064
S4 ($N = 500$)	G-BL	1.00 ± 0.00	0.490 ± 0.052	0.53 ± 0.11
	G-SM	0.37 ± 0.12	0.408 ± 0.045	0.700 ± 0.058
	GA-P-BL	1.00 ± 0.00	0.490 ± 0.052	0.53 ± 0.11
	MOEA-2	1.00 ± 0.01	0.493 ± 0.053	0.53 ± 0.11
	MOEA-3	1.00 ± 0.00	0.490 ± 0.053	0.54 ± 0.11

Table 4: Hypervolume across 200 scenarios for the two multi-objective solvers, with G-BL as a single-point reference (single-point solvers G-SM and GA-P-BL are omitted from this comparison; see § 6.4). HV comparisons between single-point and multi-point solvers are confounded by point-count asymmetry (see text).

Solver	HV	SD
MOEA-3	0.0593	0.0681
MOEA-2	0.0417	0.0398
G-BL	0.0186	0.0033




Figure 1: Statistical validation: (left) hypervolume distribution; (right) pairwise Cliff's δ effect sizes. The gray band marks negligible effects ($|\delta| < 0.147$).

Figure 1 statistically validates these solver comparisons through the hypervolume distribution and pairwise Cliff's δ effect sizes.

The following sections deepen this overall analysis from five perspectives: §6.2 examines the special trade-off of squint-aware scheduling, §6.3 investigates scale sensitivity (how multi-objective advantage varies with target density), §6.4 reveals the geometric mechanism of Pareto compression, §6.5 verifies the condition dependence of this mechanism through ablation, and §6.6 separates initialization and search-budget confounds through robustness controls.

6.2. Squint-Aware Special Regime

The squint-minimized greedy solver (G-SM) improves NESZ radiometric quality (f_3) relative to the profit-only baseline (G-BL) by selecting observation times that minimize squint angle within each visibility window. Across all 200 scenarios, this produces a large positive effect on f_3 (Cliff's $\delta = +1.0$, $p < 10^{-34}$) at a substantial aggregate cost in coverage (Cliff's $\delta = -1.0$, $p < 10^{-34}$, on f_1^* ; raw f_1 is used here because G-BL yields constant $f_1^* = 1.0$ by definition, for which the normalized-profit comparison

degenerates to a one-sample test against a fixed value). The trade-off, however, is strongly density-dependent. In sparse scenarios (S1, $N = 20$), G-SM drops to $f_1^* = 0.51 \pm 0.16$ (-49% vs. G-BL at 1.00) while improving f_3 from 0.271 ± 0.057 to 0.644 ± 0.058 ($+138\%$). In dense scenarios (S4, $N = 500$), the f_1 cost escalates: G-SM achieves $f_1^* = 0.37 \pm 0.12$ (-63%), with f_3 rising from 0.53 ± 0.11 to 0.700 ± 0.058 ($+32\%$). Figure 2 visualizes this density-dependent trade-off: the Locally Estimated Scatterplot Smoothing (LOESS) trend in panel (c) identifies a transition region around $N = 100$ where the f_1 penalty grows rapidly. The radiometric advantage narrows with density ($+138\%$ at S1 to $+32\%$ at S4) as tightly packed time windows constrain how far each acquisition can be moved toward the squint-minimizing time, while the coverage cost remains large at every density (G-SM schedules, on average, 44% of the tasks G-BL selects at S1, falling to $\sim 20\text{--}27\%$ at $N \geq 100$; these are means of the per-scenario scheduled-task-count ratios, the profit-side analogue of which appears in Table 3).




Figure 2: Effect of squint-aware scheduling: (a) f_3 scatter; (b) Δf_1^* histogram; (c) Δf_1^* vs. N with LOESS trend

6.3. Scale Sensitivity

The gap between MOEA-2 and G-BL varies systematically with scenario density (Fig. 3, Table 5). MOEA-2's coverage deficit relative to G-BL is confined to sparse densities: from $f_1^* = 0.87 \pm 0.08$ at S1 ($N = 20$) it rises to 0.985 ± 0.03 at S2 ($N = 100$), and reaches 0.997–0.999 at $N \geq 150$ (S7–S8, S3), remaining within 1% of G-BL at S4 (0.996). The f_2 improvement over G-BL follows the same contraction, falling from $+4.7\%$ (S1) to $+1.4\%$ (S2) and to $+0.8\%$ or less at $N \geq 150$. The proportion of scenarios where MOEA-2's coverage falls measurably short of G-BL ($f_1^* < 0.95$, i.e., a coverage deficit

exceeding 5%) collapses from 84% (S1) to 6% (S2) and to 0% at $N \geq 150$ (2% at S4, a single scenario). The multi-objective coverage advantage is therefore a sparse-density phenomenon: between $N = 20$ and $N = 100$ the solver transitions from actively trading off coverage for quality to operating at the greedy baseline, and beyond $N = 150$ the deficit is effectively zero.




Figure 3: Geometric quality (f_2) by scenario group and solver

Table 5: Scale sensitivity of MOEA-2 vs. G-BL. f_2 improvement is the mean of per-scenario ratios (the abstract reports the ratio of pooled means, hence +4.5% at S1).

Class	N	f_1^* (MOEA-2)	f_2 improvement	% $f_1^* < 0.95$
S1	20	0.87 ± 0.08	+4.7%	84%
S2	100	0.985 ± 0.03	+1.4%	6%
S7	150	0.998 ± 0.00	+0.8%	0%
S8	200	0.997 ± 0.01	+0.7%	0%
S3	300	0.999 ± 0.00	+0.6%	0%
S4	500	0.996 ± 0.01	+0.5%	2%

Fitting a logistic regression to the proportion of scenarios with $f_1^* < 0.95$ across the six density points ($N = 20, 100, 150, 200, 300, 500$) yields a transition midpoint $N_{50} \approx 50$ targets. This estimate is weakly constrained: only the two sparse densities ($N = 20, 100$) carry informative deficit proportions (84% and 6%), while the four denser points are all at most 2%, so the bootstrap confidence interval is wide ($N_{50} \in [30, 62]$). N_{50} should therefore be read as a coarse indicator that the transition is confined to $N \in [20, 100]$, not as a precise threshold. Beyond $N \approx 150$, at most one scenario (a single S4 case) shows a coverage deficit exceeding 5%. The transition between a measurable and a negligible multi-objective advantage therefore lies in the $N \in [20, 100]$ band, and beyond $N = 150$ the advantage is effectively absent.

6.4. Mechanism of Pareto Compression

Phenomenon 1: Front compression. Fig. 4 visualizes the Pareto front on a representative S4 scenario. The front contains 20.5 ± 9.9 MOEA-3 non-dominated solutions (MOEA-2: 1.8 ± 1.0) but is compressed: the found front spans only about 0.06 in f_1^* and 0.003 in f_2 , against a total span of $\Delta f_2 \approx 0.12$ in sparse S1 scenarios, confirming the scale-dependent saturation documented in §6.3. This compression is a property of the found front under G-BL-seeded search rather than of the full feasible trade-off surface: the squint-minimizing heuristic G-SM reaches feasible high- f_3 operating points ($f_3 \approx 0.70$ at S4, §6.2) outside the MOEA-3 surface, so the dense-regime redundancy is established near the full-coverage working point rather than over the entire feasible surface. In dense regimes, solver convergence to lower-squint geometries combined with schedule saturation limits quality differentiation regardless of solver strategy. The MOEA-2 and MOEA-3 fronts are

nearly coincident at this density, indicating that adding f_3 as a third objective produces little additional differentiation in the dense setting.




Figure 4: Pareto front on a representative S4 scenario ($N = 500$): MOEA-3 frontier with MOEA-2 overlay and baseline solver markers. Panel (a): $f_1^*-f_2$; panel (b): $f_1^*-f_3$.

The $f_1^*-f_2$ projection (panel a) shows an extremely compressed surface: even for the S4 scenario with the widest f_1^* spread, the frontier spans only $f_1^* \in [0.94, 1.00]$ and $f_2 \in [0.500, 0.503]$ — a near-vertical segment at full coverage. Overlaying the single-point baseline solutions shows that G-BL sits at the “full-coverage, moderate-quality” corner of the frontier, while G-SM occupies a qualitatively different region ($f_1^* \approx 0.37$, $f_3 \approx 0.70$) that lies outside the MOEA-3 Pareto surface, a consequence of G-SM’s aggressive squint minimization.

Phenomenon 2: Density-dependent third-objective value. The differentiation of MOEA-3 from MOEA-2 is strongly density-dependent: at S1 the physics-aware solver attains $f_3 = 0.428 \pm 0.064$ versus 0.233 ± 0.027 for MOEA-2 (+84%), at a f_2 cost of about 10% (0.540 ± 0.021 vs. 0.600 ± 0.021); at S4 the f_3 advantage collapses to +1.5% (0.535 ± 0.109 vs. 0.527 ± 0.111) with a negligible f_2 penalty. The third objective is therefore *not* redundant at sparse density, where low coverage pressure leaves the solver free to trade f_2 for f_3 ; at operational densities the advantage disappears because coverage pressure locks the selection, not because the underlying f_2 – f_3 trade-off vanishes—the within-schedule correlation actually *strengthens* with density.

The HV ranking (Table 4) differs from the coverage ranking (Table 3): among the three tabulated solvers, G-BL ranks first in tabulated f_1^* but third in HV, while the NSGA-III variants rank first and second in HV (MOEA-3, then MOEA-2), despite their lower coverage ranking. Notably, G-SM —

omitted from the table because its single-point HV is dominated by point location rather than front quality — actually attains the highest HV of any solver (0.1023) from its low-coverage, high- f_3 operating point ($f_1^* \approx 0.37$, $f_3 \approx 0.70$ at S4), so the HV ranking would invert if it were included. This is precisely the point-count confound: HV rewards an extreme single point more than a dense front near full coverage. A few sparse-density MOEA solutions exceed $f_1^* = 1.00$ (three of 200 scenarios, up to +6%), so G-BL is a near-optimal rather than a strict coverage anchor. This reversal is driven by point-count asymmetry: single-objective solvers contribute one point each to the HV reference set, while MOEA-3 contributes 53.3 ± 6.7 non-dominated points at S1 (MOEA-2: 7.3 ± 2.6), a cardinality gap that persists, less severely, at S3–S4. HV therefore conflates front quality with front cardinality, and the ranking should not be interpreted as a standalone quality judgment. The core phenomenon—front compression—is independent of HV ranking.

Mechanism: geometric coupling and lower-squint convergence.

Because θ and ψ_{sq} are not independent but both derived from the same LOS vector $\mathbf{l} = (l_x, l_y, l_z)$ (§ 3.1), the per-task f_2 and f_3 contributions share geometric structure. To separate definitional coupling from scheduler-induced correlation, we compute the analytical null distribution by treating θ and ψ_{sq} as independently and uniformly distributed within the accessible envelope ($\theta \in [15^\circ, 50^\circ]$, $|\psi_{\text{sq}}| \leq 45^\circ$, slightly wider than the C1 constraint $[18^\circ, 47^\circ]$ to provide a conservative null baseline) and evaluating the moment integrals of $f_{2,i} = \sin \theta_i \cos \psi_{\text{sq},i}$ and $f_{3,i} = \cos^3 \theta_i \cos^3 \psi_{\text{sq},i}$ in closed form (validated by Monte Carlo, $r_{\text{MC}} = -0.5120$, $N_{\text{MC}} = 2 \times 10^5$). The null correlation is

$$r_{\text{null}} = \frac{\text{Cov}(f_2, f_3)}{\sqrt{\text{Var}(f_2) \text{Var}(f_3)}} = -0.51, \quad (6)$$

where the independence of θ and ψ_{sq} under the null factorizes the moments, e.g., $E[f_2 f_3] = E[\sin \theta \cos^3 \theta] E[\cos^4 \psi_{\text{sq}}]$. The *negative* sign exposes a structural conflict, not redundancy: in the θ -only limit (fixed squint), $f_{2,i} \propto \sin \theta$ and $f_{3,i} \propto \cos^3 \theta$ are strictly anti-monotonic ($r_\theta = -0.997$), since f_2 favors high incidence while f_3 favors low incidence; in the ψ_{sq} -only limit (fixed θ), both depend on $\cos \psi_{\text{sq}}$ positively ($r_\psi = +0.997$), since low squint improves both. Under uniform geometry the θ -conflict dominates, yielding $r_{\text{null}} < 0$.

Empirical correlation and lower-squint convergence. Pooled across all scheduled tasks from the 200 scenarios, the per-task empirical correlation between f_2 and f_3 within the solver-selected schedules is negative at every

density and *strengthens* as density rises: for the full-physics MOEA-3 (variant A) it ranges from -0.61 (S1) to -0.78 (S4), and for the no-physics variant D from -0.56 (S1) to -0.78 (S4). Because the pooled statistic mixes between-schedule and within-schedule variance, we also computed per-schedule correlations: the Fisher-z mean of the per-schedule values tracks the pooled value closely (e.g. -0.64 vs. -0.61 at S1 for the full-physics variant), while individual schedules span a wide range (S1 $r \in [-0.97, +0.25]$), so the pooled value is representative but not uniform across schedules. The θ -conflict therefore persists inside the schedules rather than being reversed. All solvers select geometries shifted toward lower squint than the visibility envelope (mean $|\psi_{\text{sq}}|$ $16\text{--}24^\circ$; §3.1), which aligns f_2 with f_3 along the shared $\cos \psi_{\text{sq}}$ factor and attenuates the envelope correlation most strongly at sparse density ($-0.97 \rightarrow -0.61$ at S1); as density rises, coverage pressure forces the inclusion of tasks across a wider incidence range, widening the θ spread and strengthening the residual negative correlation (for MOEA-3, variant A, the per-schedule mean f_2 broadly falls from 0.54 at S1 to 0.49 at S4 while mean f_3 broadly rises from 0.43 at S1 to 0.54 at S4, dipping to 0.37 at S2). The third-objective redundancy is therefore an empirical convergence effect—not a mathematical identity between f_2 and f_3 but a consequence of coverage pressure that removes the trading freedom at high density while the incidence-angle conflict itself persists.

Envelope versus schedule correlation. The contrast between envelope correlation and schedule correlation is stark. Across visibility-envelope samples spanning the full time extent of each visibility window ($|\psi_{\text{sq}}| \leq 45^\circ$, computed across all visibility windows from 20 scenarios (5 per density class S1–S4, the five smallest random seeds per class), approximately 3.9×10^5 window–time-step pairs at 10 s resolution), $\text{corr}(f_2, f_3) = -0.97$, a strong negative correlation dominated by the incidence-angle conflict; the value is essentially unchanged (-0.97) if the C1 incidence interval $[18^\circ, 47^\circ]$ is additionally imposed. This envelope correlation of -0.97 differs from the factorized null $r_{\text{null}} = -0.51$ (Eq. 6) because the latter assumes independent θ and ψ_{sq} , whereas real visibility geometry couples the two angles through the LOS vector, making the θ -conflict dominate. The solver-selected schedules thus move the correlation toward zero—from -0.97 in the visibility envelope to -0.61 (S1) within the full-physics schedules—but do not reverse its sign: the incidence-angle conflict is attenuated by lower-squint convergence, not eliminated. This partial attenuation is an active effect of the solvers, not a consequence of geometric envelope constraints: envelope constraints alone

would leave the correlation at -0.97 . Ablation studies independently support this interpretation. With all physical objectives removed (variant D, § 6.5), the schedule-level correlation at S1 is $r \approx -0.56$ —closer to zero than under the full-physics variant A ($r \approx -0.61$)—and both variants converge to the same strongly negative level at high density ($r \approx -0.78$ at S4). In the sparse regime the physical objectives preserve a slightly stronger residual negative correlation by retaining solutions that traverse the f_2 – f_3 trade-off; without them, the solver maximizes coverage alone, and temporal packing incidentally favors lower-squint windows, attenuating the residual correlation further. At high density this structural packing dominates irrespective of the objective formulation, so the ablation contrast vanishes and the correlation saturates at a strongly negative level; the third objective’s value is lost not because the conflict weakens but because no formulation can escape the coverage-driven θ spread. Consistent with this scale-dependent saturation, the per-density hypervolume contracts at higher density: at S4 the pooled HV is 0.18 (MOEA-2) and 0.11 (MOEA-3) of its S1 value, i.e. the S4/S1 HV ratio is well below unity, reflecting the near-point front compression documented in §6.4 (S4 fronts span only $\Delta f_1^* \approx 0.06$) rather than a degradation of solution quality. We caution that the Pareto fronts reported here are best-known approximations obtained via NSGA-III; whether they represent the true Pareto-optimal set is unknown without exact methods, which remain intractable for $N \geq 20$ in this problem class.

6.5. Physical-Objective Ablation

To quantify the contribution of incidence-angle (θ) and squint-angle (ψ_{sq}) modeling to solver performance, we repeat the MOEA-3 experiment with four objective-function variants:

- A. Full physical (baseline)** $f_2 = \frac{1}{|S|} \sum \sin \theta_i \cdot \cos \psi_{\text{sq},i}$, $f_3 = \frac{1}{|S|} \sum \cos^3 \theta_i \cdot \cos^3 \psi_{\text{sq},i}$ (identical to the standard configuration in § 6.4).
- B. No squint** $f_2 = \frac{1}{|S|} \sum \sin \theta_i$, $f_3 = \frac{1}{|S|} \sum \cos^3 \theta_i$ (removes ψ_{sq} component).
- C. No incidence** $f_2 = \frac{1}{|S|} \sum \cos \psi_{\text{sq},i}$, $f_3 = \frac{1}{|S|} \sum \cos^3 \psi_{\text{sq},i}$ (removes θ component).
- D. No physics** $f_2 = f_3 = 1$ (constant per-task objective; solver receives no geometric information; f_2 , f_3 shown in Table 6 are post hoc diagnostics).

computed from the actual observation geometry of the resulting schedule). Variant D serves as a null-model control: it tests whether the NSGA-III solver’s performance is attributable to the physical content of f_2 and f_3 , or merely to the presence of additional objectives that structure the search.

All four variants use identical NSGA-III parameters (population = 100, generations = 200, Das-Dennis $p = 12$), the same 200 scenario files, and the same G-BL hot-start injection. Only the objective functions differ, making this a controlled ablation rather than a retuning experiment. Within-solver ablation with identical NSGA-III parameters and the same G-BL seed injection controls confounding factors sufficiently that the Wilcoxon signed-rank test is appropriate for paired comparison of objective-function variants. Each variant completes 4 scenario classes $\times$ 50 seeds = 200 runs (total: 800 runs across A–D). Degradation is reported as $\Delta\% = (A - V)/A \times 100$, where A is the full physical baseline and V is the variant; negative values indicate the variant outperforms the baseline on f_1^* .

Table 6: Per-class ablation results. Bold: $p < 0.005$ (Wilcoxon signed-rank vs. A; consistent with the Friedman–Wilcoxon framework of §5.3).

Class	V	f_1^*	f_2	f_3	f_1/n	n_{sel}	$\Delta f_1^* \%$
S1 ($N = 20$)	A	0.98 ± 0.03	0.540	0.428	5.94 ± 0.68	18.1 ± 1.7	(baseline)
	B	1.00 ± 0.03	0.539	0.405	5.79 ± 0.67	19.0 ± 1.1	− 2.0%
	C	0.90 ± 0.11	0.503	0.544	5.83 ± 0.79	17.2 ± 2.4	+ 7.5%
	D	1.01 ± 0.02	0.565	0.309	5.66 ± 0.64	19.5 ± 0.7	− 2.7%
S2 ($N = 100$)	A	1.00 ± 0.00	0.548	0.375	5.55 ± 0.38	64.0 ± 18.7	(baseline)
	B	1.00 ± 0.01	0.548	0.372	5.56 ± 0.38	64.0 ± 19.0	+0.0%
	C	0.99 ± 0.03	0.537	0.412	5.55 ± 0.39	63.3 ± 19.1	+1.2%
	D	1.00 ± 0.00	0.549	0.367	5.54 ± 0.37	64.2 ± 18.8	−0.1%
S3 ($N = 300$)	A	1.00 ± 0.00	0.528	0.453	5.63 ± 0.29	154.5 ± 39.9	(baseline)
	B	1.00 ± 0.00	0.528	0.453	5.63 ± 0.29	154.5 ± 39.9	+0.0%
	C	1.00 ± 0.00	0.525	0.461	5.62 ± 0.29	154.4 ± 39.9	+0.1%
	D	1.00 ± 0.00	0.528	0.452	5.63 ± 0.29	154.5 ± 39.9	+0.0%
S4 ($N = 500$)	A	1.00 ± 0.00	0.490	0.535	5.83 ± 0.44	142.7 ± 71.9	(baseline)
	B	1.00 ± 0.00	0.490	0.535	5.83 ± 0.44	142.7 ± 71.9	+0.0%
	C	1.00 ± 0.01	0.487	0.543	5.82 ± 0.45	142.5 ± 71.9	+0.1%
	D	1.00 ± 0.00	0.490	0.535	5.83 ± 0.44	142.7 ± 71.9	+0.0%




Figure 5: Ablation results across objective-function variants B–D and scenario classes S1–S4. Bars show degradation relative to the full-physics baseline A for normalized profit f_1^* and post hoc geometric-quality metrics f_2 and f_3 ; negative values indicate that a variant outperforms A. Asterisks mark paired Wilcoxon comparisons against A with $p < 0.005$ (Bonferroni threshold).

The ablation results (Table 6, Fig. 5) reveal that the influence of physical objective modeling is concentrated in sparse scenarios and vanishes as scenario density increases — a pattern consistent with our scale-sensitivity analysis (§ 6.3).

At S1 ($N = 20$), all three ablations differ significantly from baseline A at the manuscript’s $p < 0.005$ threshold (paired Wilcoxon). Variant C (no incidence) degrades f_1^* by +7.5% ($p = 5.8 \times 10^{-4}$) while raising post hoc f_3 from 0.428 to 0.544: removing θ leaves both objectives as functions of ψ_{sq} alone, maximized at zero squint, so C behaves like a near-G-SM strategy that trades coverage for radiometric quality. Conversely, Variant D (no physics) improves f_1^* by −2.7% ($p = 3.1 \times 10^{-6}$) by selecting about 1.4 more tasks (19.5 vs. 18.1), but its post hoc f_3 drops to 0.309 (−28% vs. 0.428 for A)

with f_2 slightly higher (0.565 vs. 0.540): with no physical guidance the solver maximizes profit and schedules more tasks, sacrificing the radiometric quality that the full-physics baseline purchases in the sparse regime. Variant B (no squint) yields a small coverage gain (-2.0% , $p = 4.9 \times 10^{-4}$) with negligible quality change. Thus physical objectives materially alter solver behavior in the sparse regime, but the direction of the coverage change depends on which component is removed.

At S2 ($N = 100$), all variants lie within 1.2% of baseline A in f_1^* , and their per-task profit (5.54–5.56) and selected-task counts (63.3–64.2) are similar: the physical-objective signal is already much weaker than at S1.

At S3–S4 ($N \geq 300$), every variant converges to within 0.2% of the baseline for f_1^* and to within 0.6% for f_2 ; the per-task quality metrics (5.62–5.83) and selected-task counts agree to within 0.2. This numerical invariance is the direct signature of the scale-dependent saturation discussed in § 6.3: once the density of candidate observation windows saturates the schedule, the objective function only penalizes infeasible solutions rather than differentiating among feasible ones. In this regime, the NSGA-III solver acts as a constraint-filtered profit maximizer regardless of whether f_2 and f_3 carry geometric meaning (A) or serve as constant placeholders (D).

Taken together, the physical quality objectives add value only in the sparse regime: at S1 the full-physics A attains +38% higher f_3 than the no-physics D (0.428 vs. 0.309) at a 2.7% coverage cost, a signal already weakened by S2 ($N = 100$) and numerically absent at $N \geq 300$.

6.6. Robustness and Calibration

The mechanism above must be separated from initialization and search-budget effects. We therefore use four control experiments.

Hot-start dependence.. On 10 S1, 4 S2, 5 S3, and 5 S4 scenarios (3 seeds each), random-initialized MOEA-2 yields f_1^* deficits of approximately -0.237 ± 0.114 , -0.480 ± 0.063 , -0.622 ± 0.062 , and -0.461 ± 0.092 (mean $\pm$ SD across runs) relative to hot-start MOEA-2, respectively. G-BL seeding therefore materially improves coverage convergence at every tested density. The deficit peaks at the dense S3 class (-0.62 at $N = 300$), is smallest at the sparsest S1 (-0.24 at $N = 20$), and moderates at the densest S4 (-0.46 at $N = 500$), so high target density does not by itself make random initialization sufficient. This constrains the coverage-ceiling claim: proximity to G-BL is initialization-dependent. By contrast, random-init f_2 is essentially

unchanged at every density (0.551 vs. 0.550 at S3, 0.551 vs. 0.551 at S4), so the lower-squint regime is not an artifact of the anchor; the coverage deficit is confined to f_1^* . The coverage ceiling is therefore a G-BL initialization effect rather than an artifact of the MOEA framework. Two caveats qualify this control: the deficit sample is small (10/4/5/5 scenarios per class $\times$ 3 seeds, and 30 runs for the variant-D control, 15 per class), and the comparison is expressed on the G-BL-normalized f_1^* scale, on which the greedy anchor itself reads exactly 1.0 by construction; both argue for reading the deficit magnitudes as indicative of the initialization effect rather than as precise coverage bounds.

Perturbation sensitivity.. A sweep over $\sigma \in \{0.1, 0.3, 0.5, 0.7\}$ on MOEA-3 (10 S3 scenarios per value, fixed solver seed) leaves mean f_2 and f_3 essentially unchanged (mean $f_2 \approx 0.535$ and mean $f_3 \approx 0.435$ at every level, with level-to-level variation of at most $\sim 2 \times 10^{-3}$; MOEA-2 shows the same flatness at $f_2 \approx 0.537$ and $f_3 \approx 0.429$), and coverage stays flat at 200/300 selected tasks. The per-task f_2 – f_3 coupling reported in § 6.4 is therefore a property of the schedule geometry rather than of the perturbation magnitude, and redundancy persists across $\sigma = 0.1$ – 0.7 rather than being a byproduct of the $\sigma = 0.5$ perturbation level.

Random-search calibration.. On five S1 scenarios, 5,000 random task selections per scenario produce mean $f_1^* = 0.527 \pm 0.003$ and a 90th percentile of 0.930 ± 0.008 ; the best sample reaches 1.000 — the G-BL profit, so random sampling can at most replicate the greedy baseline, never exceed it. High coverage under random sampling therefore occurs only by chance; the MOEA’s gains reflect structured search.

Random-init control.. A random-init variant-D control (15 S3 and 15 S4 runs, no G-BL hot start) also converges to low squint, with $f_2 = 0.527 \pm 0.013$ at S3 and 0.520 ± 0.009 at S4. At S3 this matches the hot-started D solver ($f_2 = 0.527$), while at the densest S4 it exceeds it (0.520 vs. 0.490), so the lower-squint convergence is objective-intrinsic rather than a hot-start residue: the G-BL seed inserts tasks at their earliest feasible time and thus biases hot-started runs toward the higher squint of the greedy schedule, most visibly at high density. The A–D invariance itself already holds at standard budget (Table 6: $\Delta f_1^* \leq 0.1\%$ at S3–S4), a physically forced consequence of coverage pressure locking the selection at high density (§ 6.4); the convergence is

therefore not a budget artifact. These controls collectively favor the lower-squint convergence explanation over hot-start residue, noise-level sensitivity, or random chance, while not ruling out all algorithm-specific effects.

7. Discussion

7.1. Summary of Findings

We define a solver as *Pareto-sufficient* at scale N if no alternative solver in the comparison set achieves both higher f_1^* and higher f_2 (or f_3) with statistical significance ($p < 0.005$, Cliff’s $|\delta| > 0.33$). This operational criterion separates cases where multi-objective optimization yields a measurable quality–coverage advantage from cases where it does not. Under the window-corrected data, no solver attains Pareto-sufficiency even at the sparsest density: at $N = 20$ the three dimensions form a genuine trade-off surface — G-BL retains $f_1^* = 1.00$ at a moderate f_3 (0.271), MOEA-2 reaches the highest f_2 (0.600) at the lowest f_3 (0.233), and MOEA-3 trades f_2 (−10% vs. MOEA-2) for the highest f_3 (0.428, +84% vs. MOEA-2) — so no solver dominates on more than one dimension, whereas at $N \geq 300$ all solvers converge to the G-BL anchor and the greedy baseline is operationally sufficient. Our experiment across 200 systematically varied scenarios yields five interconnected findings: **(1)** the quality advantage of multi-objective optimization saturates with target density (+4.7% in f_2 at $N = 20$ falling to +0.5% at $N = 500$); **(2)** the third objective is density-dependent rather than redundant: per-task f_2 – f_3 correlation within schedules stays negative and *strengthens* with density (≈ -0.61 to -0.78 for the full-physics solver), and MOEA-3 separates from MOEA-2 in f_3 at sparse density (+84% at $N = 20$) with the advantage shrinking to $\approx +1.5$ –2% at $N \geq 300$, confirming the coverage-pressure mechanism detailed in § 6.4; **(3)** G-SM occupies a distinct high-radiometric-quality regime at substantial density-dependent coverage cost; **(4)** solver outcomes become invariant to the choice of physical objectives at $N \geq 300$ (ablation study); and **(5)** coverage convergence toward the G-BL anchor is conditional on the hot start, not a structural property. The first four reflect geometric properties of the single-pass observation geometry (finding (2) being an empirical convergence effect under those properties, § 6.4); the fifth reflects solver initialization.

The Friedman test on pooled HV confirms overall differences among the five solvers across all 200 scenarios ($\chi^2 = 552.91$, $p \approx 0$; G-BL and GA-P-BL tie by construction, so the statistic is conservative toward fewer distinct

rankings). Pairwise Wilcoxon signed-rank tests use the Bonferroni threshold $\alpha = 0.005$; interpretation relies on Cliff’s δ given the large sample size. MOEA-2 and MOEA-3 differ on pooled HV with $\delta = -0.58$ ($p < 10^{-14}$), but this contrast blends the sparse-regime divergence with the dense-regime convergence and is inflated by the HV point-count asymmetry (§ 6.4); contrasts involving G-SM are large because it occupies an extreme radiometric-quality regime.

Independently of the multi-objective question, GA-P-BL output is numerically identical to G-BL, because its never-worse-than-baseline guard (Sec. 4.4) returns the greedy schedule on the single profit objective; the present claims therefore concern the marginal value of adding physical objectives rather than profit-oriented search alone.

7.2. Operational Implications

The scale dependence has direct consequences for mission planning. In sparse regimes ($N \leq 100$), MOEA-2 achieves a +4.5% geometric-quality improvement at S1 at a 13% coverage cost ($f_1^* = 0.87$ at $N = 20$ relative to G-BL). Whether this trade-off is acceptable depends on the application: for maritime surveillance tasks where a few-percent resolution improvement (~ 0.3 m at C-band: $\rho_{\text{rg}} \approx 5.3$ m at $\theta = 30^\circ$) may affect whether adjacent targets within the same pixel are resolved or remain merged at baseline resolution (Curlander and McDonough, 1991), the coverage cost may be justified; for area-monitoring missions, the greedy baseline remains preferable. In dense regimes ($N \geq 300$), the multi-objective advantage falls below 2% improvement in f_2 , while G-SM offers the highest radiometric quality (+32% f_3 at S4) at a 63% coverage cost; this gain comes from the one-line squint-minimization heuristic rather than from multi-objective search. The logistic midpoint $N_{50} \approx 50$ (§ 6.3) provides a coarse indicator estimate between these regimes. Our results suggest a conditional operational guideline: for sparse observation requests ($N \lesssim 100$), MOEA-2 may provide geometric quality gains that outweigh the coverage penalty in specific contexts; for dense monitoring campaigns ($N \gtrsim 300$), greedy methods (G-SM for radiometric priority, G-BL for coverage priority) attain the tested family’s coverage ceiling and may be preferable. Indeed, the MOEA-2 coverage gap relative to the coverage-only anchor is at most 2% at $N \geq 300$ (Table 5); since GA-P-BL output is identical to G-BL, no separate gap exists.

7.3. Limitations

Our study is bounded by scope decisions that define directions for future work. The 3-objective redundancy reported above is specific to the single-pass, single-satellite Stripmap setting: it reflects solver convergence to lower-squint geometries that attenuate the envelope-level f_2 – f_3 correlation while the residual within-schedule θ -conflict persists and strengthens with density (§ 6.4, Eq. 6), not a mathematical identity between f_2 and f_3 . In multi-pass or multi-satellite settings with greater angular diversity, the θ -conflict would re-emerge and the third objective may become discriminative. Similarly, ScanSAR, Spotlight, and TOPS modes introduce different NESZ and resolution dependencies beyond Stripmap, and the scenario design uses a fixed orbit (693 km circular LEO) without exploring sensitivity to orbital parameters. The hot-start control experiment (§ 6.6) indicates that G-BL initialization materially improves MOEA convergence; the coverage convergence reported here is conditional on this initialization, and full-scale random-init experiments are needed to establish the unconditional behavior. The hot-start dependence implies that the coverage-ceiling result reflects the combination of G-BL prior knowledge with MOEA search; whether the MOEA can independently discover high-coverage solutions without greedy seeding remains an open question. Similarly, the NSGA-III crowding-distance explanation for the S1 variant-D reversal (§ 6.5) awaits verification with alternative crowding operators. Two further scope notes bound the operational reading: the energy and memory budgets (C3–C4) are deliberately set far below their caps, so the saturation result holds in the resource-unconstrained regime and may not transfer when energy or downlink capacity is the binding constraint; and the dwell time is fixed at 30 s, removing the dwell-versus-SNR trade-off that an agile scheduler could otherwise exploit. The reference-point counts of MOEA-2 ($H = 13$) and MOEA-3 ($H = 91$) also differ by construction, so the sparse-regime third-objective advantage is partly entangled with reference-direction density. Finally, we do not benchmark against learning-based schedulers (e.g., Chen et al. (2025)), which represent an alternative methodological family; a systematic comparison of physics-aware heuristic methods against data-driven approaches is left to future work.

8. Conclusion

We described a physics-aware multi-objective optimization approach for agile SAR satellite scheduling that integrates NESZ radiometric quality and

geometric resolution as explicit optimization objectives alongside coverage profit. Five solvers (spanning greedy heuristics, single-objective evolutionary search, and NSGA-III with two and three objectives) were compared on 200 systematically varied scenarios across four density classes. The central finding is a scale-dependent empirical saturation of the marginal benefit of physical objectives in G-BL-seeded MOEA formulations for single-pass, single-satellite Stripmap SAR operations. At sparse density the three objectives form a genuine trade-off surface (MOEA-2 raises f_2 by +4.7% and MOEA-3 raises f_3 by +58% at modest coverage cost), while at operational densities the marginal quality gains become negligible and all solvers converge to the G-BL anchor. The saturation arises from solver convergence to lower-squint regimes under which the geometric coupling aligns f_2 and f_3 along the shared $\cos \psi_{\text{sq}}$ factor, rendering the third objective redundant only at operational densities, not at sparse density. Coverage convergence toward the G-BL anchor is conditional on the hot start; the geometric coupling is the anchor-independent result (§ 6.4). GA-P-BL output is identical to the G-BL anchor, because its never-worse-than-baseline guard always triggers on the single profit objective, so this result delimits the benefit of adding physical objectives. This saturation boundary applies to the evaluated target-density range ($N = 20\text{--}500$). For mission planning, the operational implication is conditional: multi-objective optimization provides measurable benefit only when observation requests are sparse; at dense monitoring cadences, greedy methods attain the tested family’s coverage ceiling. These findings establish an empirical saturation boundary for what physics-aware multi-objective scheduling can contribute to single-satellite SAR operations, and provide a baseline against which multi-satellite, multi-pass, and multi-mode extensions can be measured.

Declaration of generative AI and AI-assisted technologies in the manuscript preparation process

During the preparation of this work the author used AI-assisted tools (Claude, Codex) for manuscript preparation and revision assistance only. The author reviewed and edited all AI-assisted output and takes full responsibility for the content of this article.

Acknowledgments

The author thanks Prof. LIU Fanming (Harbin Engineering University) for guidance and support throughout this research. Prof. Liu is a co-author of Wu et al. (2026); his contribution to the present work was limited to advisory discussions on research direction and did not include experimental design, data analysis, or manuscript preparation. This work was supported by no specific external grant; it was performed as part of the author’s doctoral studies at Harbin Engineering University.

CRedit Authorship Contribution Statement

Liu Shuhao: Conceptualization, Methodology, Software, Validation, Formal Analysis, Investigation, Resources, Data Curation, Writing – Original Draft, Writing – Review & Editing.

Data Availability Statement

Solver code, scenario generators, fixed seeds and configurations, processed JSON/CSV results, analysis scripts, and figures are available from the corresponding author. The companion repository is publicly available at <https://github.com/liushuhao/sar-quality-aware-scheduling>. Generated binary scenario files and per-run pickle objects are not tracked; they can be regenerated using the supplied scripts.

Declaration of Competing Interest

The author declares that he has no known competing financial interests or personal relationships that could have appeared to influence the work reported in this paper.

Funding

This research received no specific grant from any funding agency in the public, commercial, or not-for-profit sectors. The work was performed as part of the author’s doctoral studies at Harbin Engineering University, College of Intelligent Systems Science and Engineering.

References

- Bianchessi, N., Cordeau, J.F., Desrosiers, J., Laporte, G., Raymond, V., 2007. A heuristic for the multi-satellite, multi-orbit and multi-user management of Earth observation satellites. *European Journal of Operational Research* 177, 750–762. doi:10.1016/j.ejor.2005.12.026.
- Caleiras, M., Moniz, S., Nascimento, P.J., 2026. A constraint programming model for the super-agile earth observation satellite imaging scheduling problem, in: *Proc. Int. Conf. on Operations Research and Enterprise Systems (ICORES)*. arXiv:2601.11967. preprint; ICORES 2026 proceedings, to appear.
- Chang, Z., Zhou, Z., 2025. Three multi-objective memetic algorithms for observation scheduling problem of active-imaging agile earth observation satellites. *Annals of Operations Research* 346, 861–893. doi:10.1007/s10479-024-06156-5.
- Chang, Z., Zhou, Z., Li, R., 2022. Observation scheduling for a state-of-the-art sar earth observation satellite: Two adaptive multi-objective evolutionary algorithms. *Computers & Industrial Engineering* 169. doi:10.1016/j.cie.2022.108252.
- Chen, M., Du, Y., Tang, K., Xing, L., Chen, Y., Chen, Y., 2024. Learning to construct a solution for the agile satellite scheduling problem with time-dependent transition times. *IEEE Transactions on Systems, Man, and Cybernetics: Systems* 54, 5991–6002. doi:10.1109/tsmc.2024.3411640.
- Chen, M., Wei, L., Chun, J., He, L., Xiang, S., Xing, L., Chen, Y., 2025. A neural priority model for agile earth observation satellite scheduling using deep reinforcement learning. *Applied Soft Computing* 174. doi:10.1016/j.asoc.2025.112984.
- Chen, X., Reinelt, G., Dai, G., Spitz, A., 2019. A mixed integer linear programming model for multi-satellite scheduling. *European Journal of Operational Research* 275, 419–434. doi:10.1016/j.ejor.2018.11.058.
- Cumming, I., Wong, F., 2005. *Digital Processing of Synthetic Aperture Radar Data: Algorithms and Implementation*. Artech House.

- Curlander, J., McDonough, R., 1991. Synthetic Aperture Radar: Systems and Signal Processing. Wiley.
- Deb, K., Jain, H., 2014. An evolutionary many-objective optimization algorithm using reference-point-based nondominated sorting approach, part i: Solving problems with box constraints. *IEEE Transactions on Evolutionary Computation* 18, 577–601. doi:10.1109/TEVC.2013.2281535.
- Ferrari, B., Cordeau, J.F., Delorme, M., Iori, M., Orosei, R., 2025. Satellite scheduling problems: A survey of applications in earth and outer space observation. *Computers & Operations Research* 173. doi:10.1016/j.cor.2024.106875.
- Herrmann, A., Schaub, H., 2023. Reinforcement learning for the agile earth-observing satellite scheduling problem. *IEEE Transactions on Aerospace and Electronic Systems* , 1–13doi:10.1109/taes.2023.3251307.
- Jacquet, A., Infantes, G., Benazera, E., Roussel, S., Meuleau, N., Baudoui, V., Guerra, J., 2024. Earth observation satellite scheduling with graph neural networks and monte carlo tree search, in: SPAICE2024 / IWPSS 2025. arXiv:2408.15041.
- Kramer, E., Miller, D., 2024. Visibility metric for planning synthetic aperture radar observations in challenging terrain. *IEEE Transactions on Geoscience and Remote Sensing* 62, 1–14. doi:10.1109/tgrs.2024.3487050.
- Kramer, E., Miller, D., 2025. Scheduling observations of visibility stressing terrain with SAR satellite constellations. *Acta Astronautica* 235, 405–420. doi:10.1016/j.actaastro.2025.05.017.
- Lemaître, M., Verfaillie, G., Jouhaud, F., Lachiver, J.M., Bataille, N., 2002. Selecting and scheduling observations of agile satellites. *Aerospace Science and Technology* 6, 367–381. doi:10.1016/S1270-9638(02)01173-2.
- Li, L., Feng, Q., Li, K., Liang, Y., 2026. Efficient generation and quality screening of visible windows for regional SAR reconnaissance. arXiv:2604.19539. preprint.
- Liu, Z., Xiong, W., Jia, Z., Han, C., 2024. Two-stage deep reinforcement learning method for agile optical satellite scheduling problem. *Complex & Intelligent Systems* 11, 35. doi:10.1007/s40747-024-01667-x.

- Moreira, A., Krieger, G., Villano, M., Younis, M., Prats-Iraola, P., Zink, M., 2026. Spaceborne synthetic aperture radar: Future technologies and mission concepts. *Proceedings of the IEEE* 114, 183–219. doi:10.1109/jproc.2025.3621586.
- Peng, G., Dewil, R., Verbeeck, C., Gunawan, A., 2019. An orienteering problem with time-dependent profits and travel times. *Computers & Operations Research* 111, 84–98. doi:10.1016/j.cor.2019.05.030.
- Wang, F., Chen, Y., He, L., Chen, J., 2026. A lagrangian relaxation-based heuristic algorithm for multiple agile earth observation satellite scheduling with time-dependent constraint. *European Journal of Operational Research* 332, 84–100. doi:10.1016/j.ejor.2025.12.047.
- Wang, J., Demeulemeester, E., Hu, X., Qiu, D., Liu, J., 2019. Exact and heuristic scheduling algorithms for multiple earth observation satellites under uncertainties of clouds. *IEEE Systems Journal* 13, 3556–3567. doi:10.1109/JSYST.2018.2874223.
- Wei, L., Cui, Y., Chen, M., Wan, Q., 2025. Multi-objective neural policy approach for agile earth satellite scheduling problem considering image quality. *Swarm and Evolutionary Computation* 94. doi:10.1016/j.swevo.2025.101857.
- Wu, K., Zhang, D., Chen, Z., Chen, J., Shao, X., 2019. Multi-type multi-objective imaging scheduling method based on improved NSGA-III for satellite formation system. *Advances in Space Research* 63, 2551–2565. doi:10.1016/j.asr.2019.01.006.
- Wu, W., Wu, X., Liu, F., Zhu, J., 2026. Improved whale optimization algorithm for increasing the task planning efficiency of agile earth observation satellites under a high target density. *Acta Astronautica* 238, 1–11. doi:10.1016/j.actaastro.2025.09.005.
- Yao, F., Chen, Y., Wang, L., Chang, Z., 2024. Bilevel evolutionary algorithm for large-scale multiobjective task scheduling in multi-agile earth observation satellite systems. *IEEE Transactions on Systems, Man, and Cybernetics: Systems* 54, 3512–3524. doi:10.1109/tsmc.2024.3359265.

Zhou, G., Jin, Z., Liu, D., 2025. Deep reinforcement learning-based two-phase hybrid optimization for scheduling agile earth observation satellites. *Remote Sensing* 17. doi:10.3390/rs17121972.